



```
[11]: plt.legend(legend, loc='best')
```



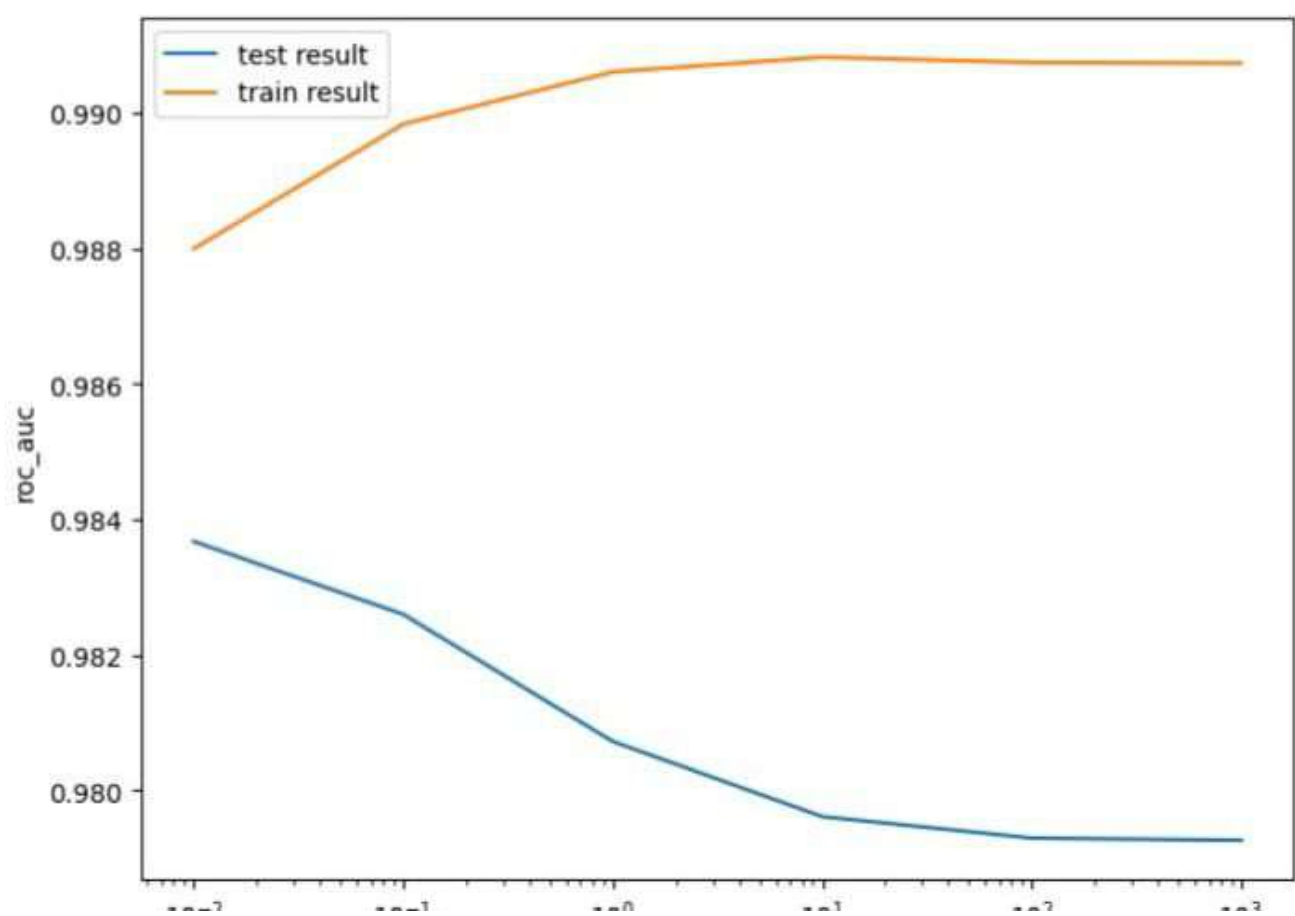
File Edit View Run Kernel Settings Help

draw_roc(y_test, y_test_pred_proba)




File Edit View Run Kernel Settings Help

🔍 + ✂️ 📄 📄 ▶️ ⏏️ 🔄 ⏏️ Code ▾




File Edit View Run Kernel Settings Help

📄 + ✂ 📄 ▶ ■ ↺ ⏏ Code ▾




File Edit View Run Kernel Settings Help

Code




File Edit View Run Kernel Settings Help

draw_roc(y_test, y_test_pred_proba)




File Edit View Run Kernel Settings Help

📄 + ✂ 📄 ▶ ■ ↺ ⏏ Code ▾




File Edit View Run Kernel Settings Help

📁 + ✂ 📄 📄 ▶ ■ 🔄 ⏏ Code ▾




File Edit View Run Kernel Settings Help

📄 + ✂ 📄 ▶ ■ ↺ ⏭ Code ▾




File Edit View Run Kernel Settings Help

⏏ + ✂ 📄 📄 ▶ ■ ↺ ⏏ Code ▾

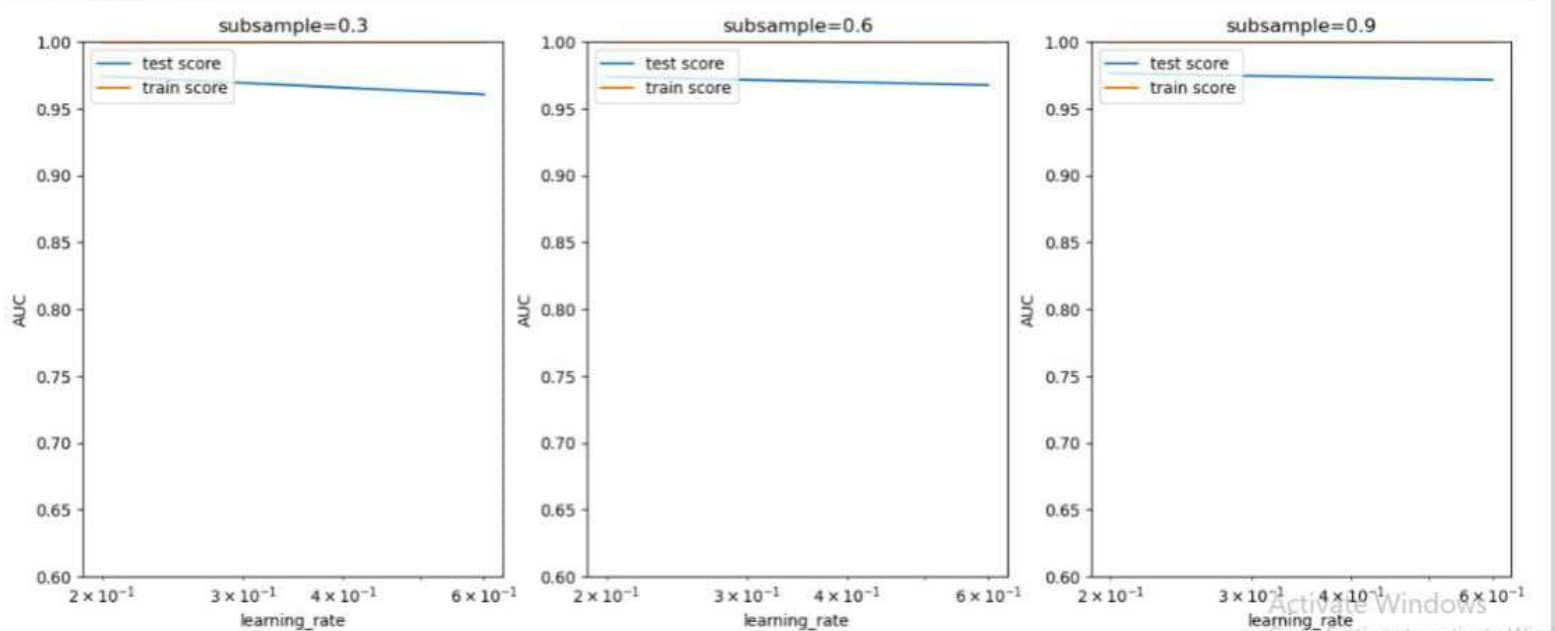


dit View Run Kernel Settings Help

✂ 📄 🗑 ▶ ⏏ 🔁 ⏩ Code ▾





We can see that we have very good ROC on the test set 0.97, which is almost close to 1.



```
train_roc(y_train, y_train_pred_proba)
```

Receiver operating characteristic example





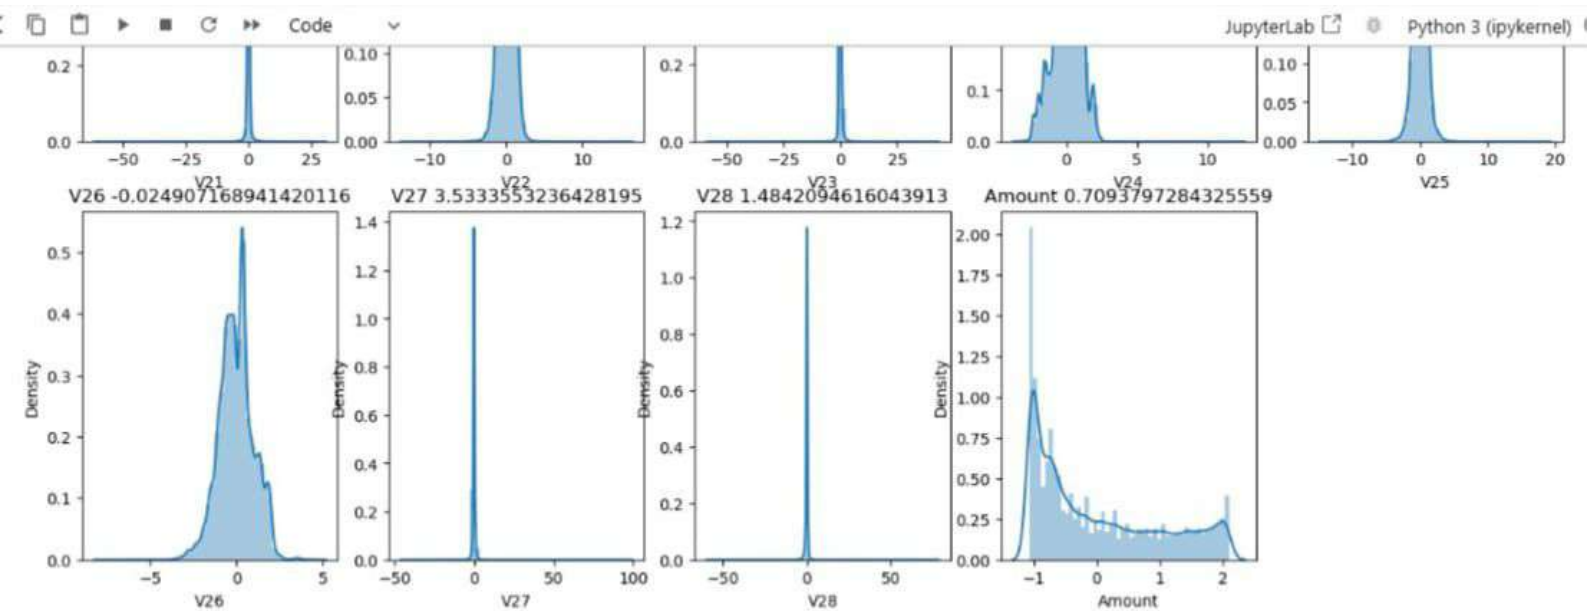
On Github, the HTML representation is unable to render, please try loading this page with noviewer.org.

[30]:

results of grid search CV
cv_results = pd.DataFrame(model_cv.cv_results_)
cv_results

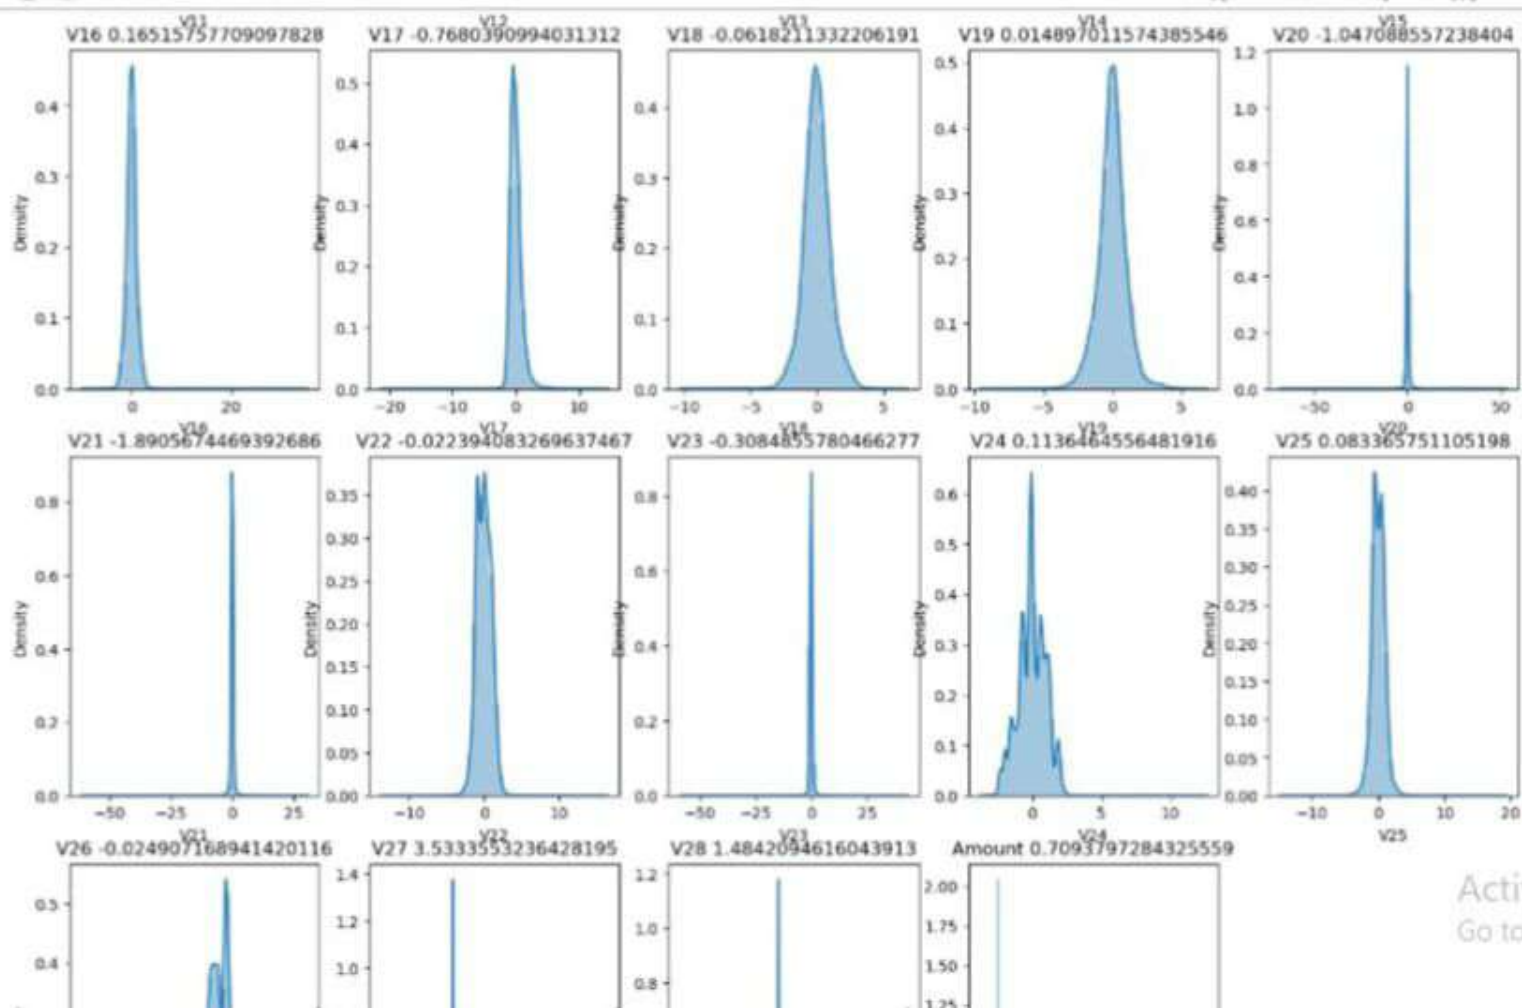
[30]:

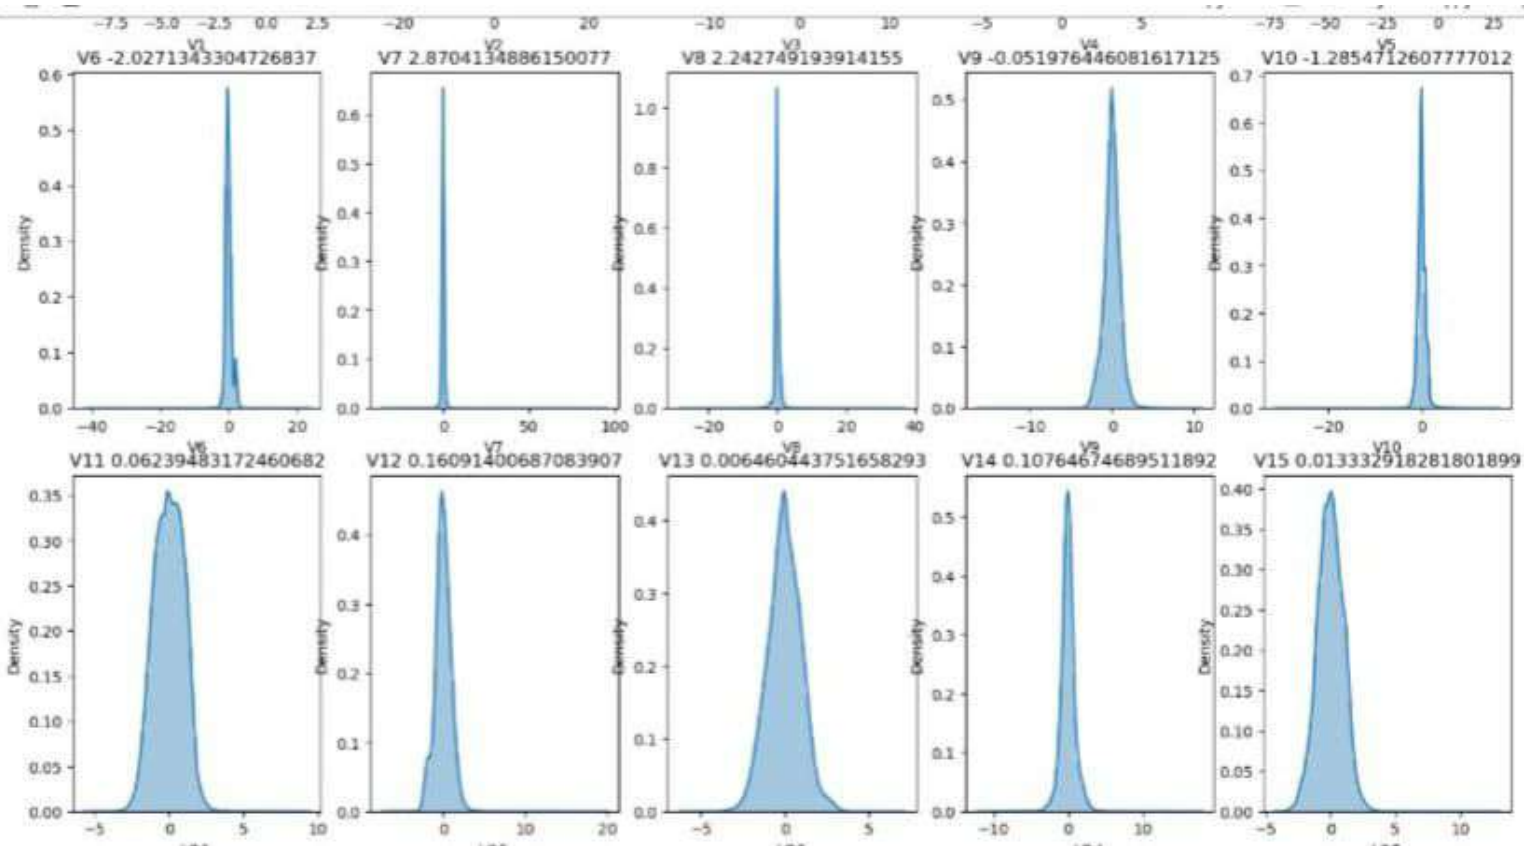
in_test_score	std_test_score	rank_test_score	split0_train_score	split1_train_score	split2_train_score	split3_train_score	split4_train_score	mean_train_score	std_train_score
0.983719	0.008479	1	0.984043	0.984587	0.988474	0.985596	0.983075	0.985155	0.001849
0.981416	0.010893	2	0.982402	0.983785	0.987917	0.984018	0.981187	0.983862	0.002270
0.980484	0.011635	3	0.981722	0.983322	0.987492	0.983305	0.980489	0.983266	0.002365
0.980375	0.011715	4	0.981632	0.983262	0.987435	0.983216	0.980404	0.983190	0.002375
0.980365	0.011722	5	0.981625	0.983256	0.987429	0.983207	0.980396	0.983182	0.002376
0.980363	0.011723	6	0.981623	0.983256	0.987428	0.983206	0.980395	0.983182	0.002376



Now we can see that all the variables are normally distributed after the transformation.

Activate Windows





File Edit View Run Kernel Settings Help

📁 + ✂ 📄 📄 ▶ ■ ↺ ⏮ Code ▼

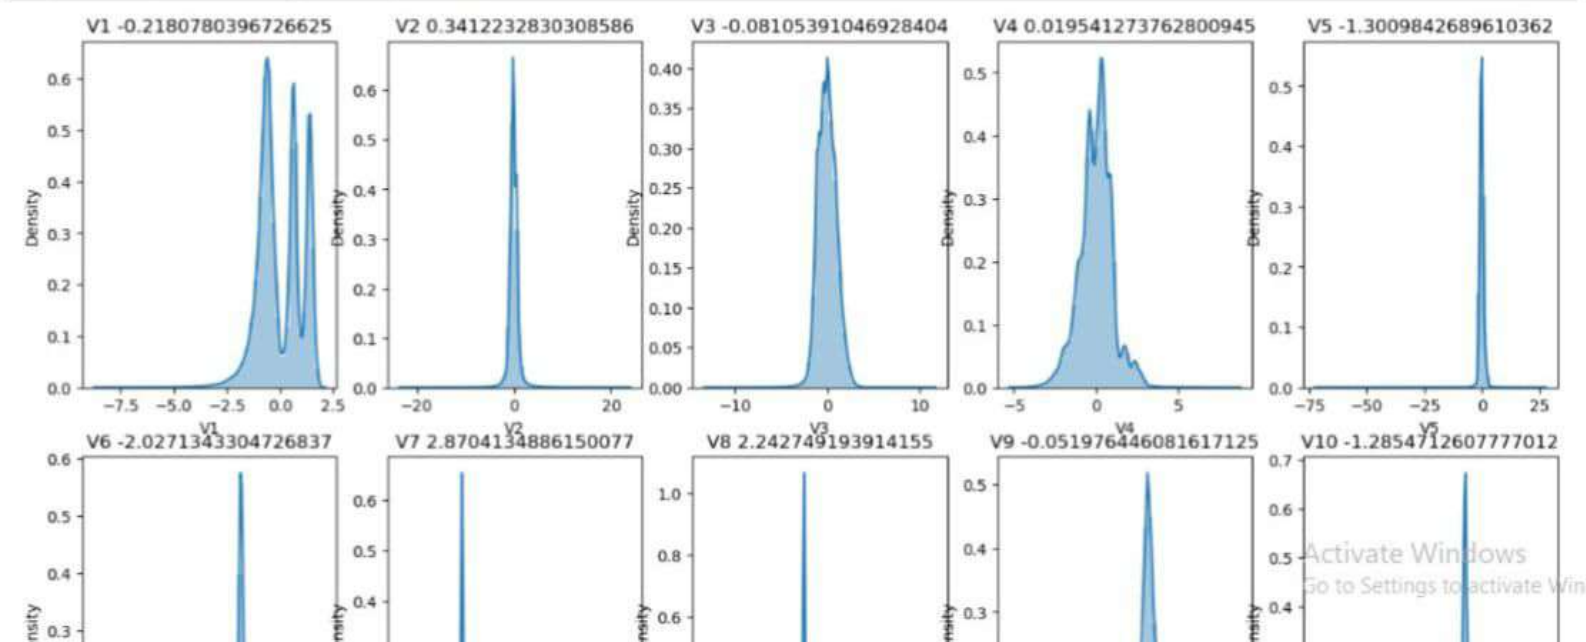
Checking the Skewness

```
[22]: # Listing the columns  
cols = X_train.columns  
cols
```

```
[22]: Index(['V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10', 'V11',  
        'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20', 'V21',  
        'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount'],  
        dtype='object')
```

```
[23]: # Plotting the distribution of the variables (skewness) of all the columns  
k=0  
plt.figure(figsize=(17,28))  
for col in cols :  
    k=k+1  
    plt.subplot(6, 5,k)  
    sns.distplot(X_train[col])  
    plt.title(col+' '+str(X_train[col].skew()))
```

```
plt.subplot(6, 5, k)
sns.distplot(X_train[col])
plt.title(col+' '+str(X_train[col].skew()))
```



Scaling the test set

We don't fit scaler on the test set. We only transform the test set.

```
[21]: # Transform the test set
X_test['Amount'] = scaler.transform(X_test[['Amount']])
X_test.head()
```


[21]:		V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	...	V20	V21	V22	V23	
	49089	1.229452	-0.235478	-0.627166	0.419877	1.797014	4.069574	-0.896223	1.036103	0.745991	-0.147304	...	-0.057922	-0.170060	-0.288750	-0.130270	1.025
	154704	2.016893	-0.088751	-2.989257	-0.142575	2.675427	3.332289	-0.652336	0.752811	1.962566	-1.025024	...	-0.147619	-0.184153	-0.089661	0.087188	0.570
	67247	0.535093	-1.469185	0.868279	0.385462	-1.439135	0.368118	-0.499370	0.303698	1.042073	-0.437209	...	0.437685	0.028010	-0.384708	-0.128376	0.286
	251657	2.128486	-0.117215	-1.513910	0.166456	0.359070	-0.540072	0.116023	-0.216140	0.680314	0.079977	...	-0.227278	-0.357993	-0.905085	0.223474	-1.075
	201903	0.558593	1.587908	-2.368767	5.124413	2.171788	-0.500419	1.059829	-0.254233	-1.959060	0.948915	...	0.249457	-0.035049	0.271455	0.381606	0.332

5 rows x 29 columns

Edit View Run Kernel Settings Help

✂ 📄 📄 ▶ ■ 🔄 ▶▶ Code ▾


```
ax = sns.violinplot(data_non_fraud['time'], label='non-fraudulent', color='blue')  
ax.set(xlabel='Transction Amount')  
plt.show()
```



File Edit View Run Kernel Settings Help


Code

previews()



File Edit View Run Kernel Settings Help

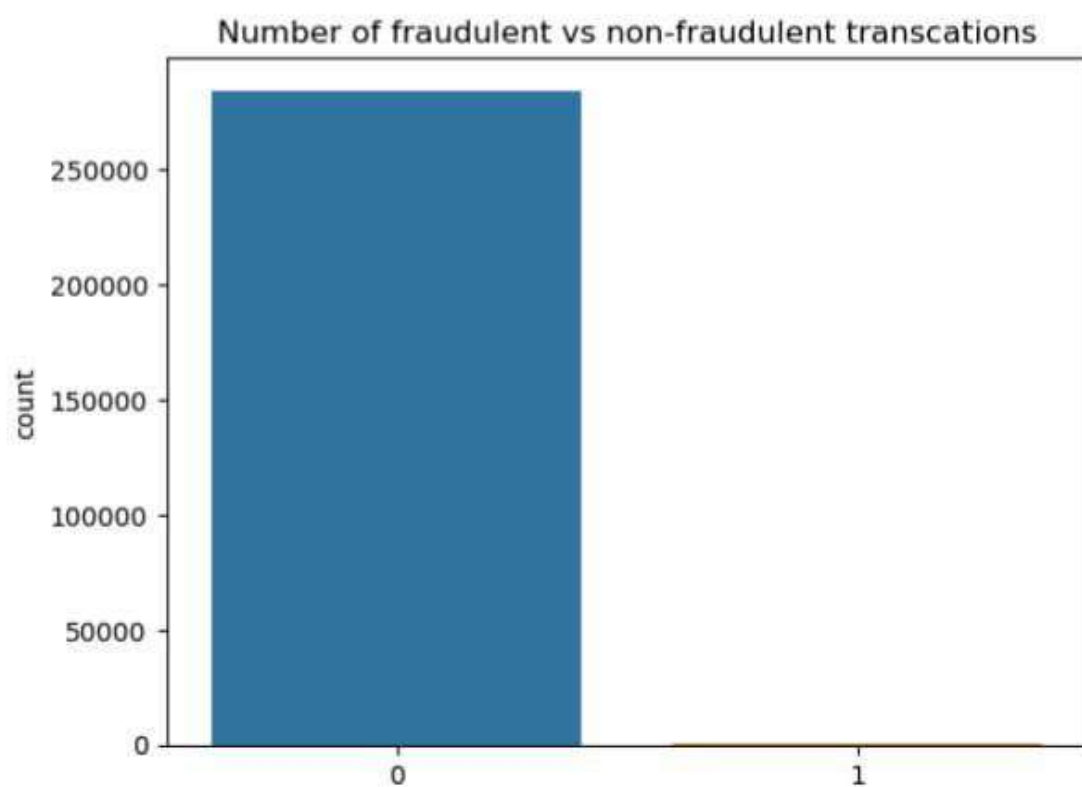
Code



File Edit View Run Kernel Settings Help

📄 + ✂ 📄 ▶ ■ ↺ ⏏ Code ▾

```
[10]: # Bar plot for the number of fraudulent vs non-fraudulent transctions
sns.countplot(x='Class', data=df)
plt.title("Number of fraudulent vs non-fraudulent transctions")
plt.show()
```



Checking the distribution of the classes

```
[7]: classes = df['Class'].value_counts()  
classes
```

```
[7]: Class  
0    284315  
1       492  
Name: count, dtype: int64
```

```
[8]: normal_share = round((classes[0]/df['Class'].count()*100),2)  
normal_share
```

```
[8]: 99.83
```

```
[9]: fraud_share = round((classes[1]/df['Class'].count()*100),2)  
fraud_share
```

```
[9]: 0.17
```

We can see that there is only 0.17% frauds. We will take care of the class imbalance later.

```
[10]: # Bar plot for the number of fraudulent vs non-fraudulent transctions  
sns.countplot(x='Class', data=df)  
plt.title('Number of fraudulent vs non-fraudulent transctions')  
plt.show()
```

File Edit View Run Kernel Settings Help

Save + Copy Paste Run Refresh Code

JupyterLab

```
[6]: # Cheking percent of missing values in columns
df_missing_columns = (round(((df.isnull().sum()/len(df.index))*100),2).to_frame('null')).sort_values('null', ascending=False)
df_missing_columns
```

```
[6]:
```

	null
Time	0.0
V16	0.0
Amount	0.0
V28	0.0
V27	0.0
V26	0.0
V25	0.0
V24	0.0
V23	0.0
V22	0.0
V21	0.0
V20	0.0
V19	0.0

```
File Edit View Run Kernel Settings Help
+ ✂ 📄 📄 ▶ ■ 🔍 ⏪ Code
JupyterLab Python 3 (ipykernel)

[5]: df.describe()

[5]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	...	2.84
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15	9.604066e-16	1.487313e-15	-5.556467e-16	1.213481e-16	-2.406331e-15	...	1.6
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00	1.098632e+00	...	7.3
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01	-1.343407e+01	...	-3.4
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01	-6.430976e-01	...	-2.2
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02	-5.142873e-02	...	-2.9
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01	5.971390e-01	...	1.8
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01	1.559499e+01	...	2.7

8 rows × 31 columns

File Edit View Run Kernel Settings Help

+ ✂ 📄 📄 ▶ ■ ↺ ⏪ Code ▾

[3]: (284807, 31)

[4]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column  Non-Null Count  Dtype  
---  --
 0   Time    284807 non-null float64
 1   V1      284807 non-null float64
 2   V2      284807 non-null float64
 3   V3      284807 non-null float64
 4   V4      284807 non-null float64
 5   V5      284807 non-null float64
 6   V6      284807 non-null float64
 7   V7      284807 non-null float64
 8   V8      284807 non-null float64
 9   V9      284807 non-null float64
10  V10     284807 non-null float64
11  V11     284807 non-null float64
12  V12     284807 non-null float64
13  V13     284807 non-null float64
14  V14     284807 non-null float64
15  V15     284807 non-null float64
16  V16     284807 non-null float64
17  V17     284807 non-null float64
18  V18     284807 non-null float64
19  V19     284807 non-null float64
20  V20     284807 non-null float64
```

```
[21]: # Read csv File from locally stored file :

file_path = r"C:\Users\Admin\Desktop\creditcard\creditcard.csv"
df = pd.read_csv(file_path)
df.head()
```

[21]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	-0.181
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	0.121
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	-0.131
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	-0.221
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	0.501

5 rows × 31 columns



```
[3]: df.shape
```

[3]: (284807, 31)